

<https://slurm.schedmd.com/sacct.html#lbaF>

column	description
JobID	Number of the job or jobstep. It is in the form: job.jobstep.
JobIDRaw	If there is an array of jobs, gets the current JobID and not the Id of the list of jobs
Cluster	Name of the computing cluster the job was run on
Partition	The partition where the job was run (production, pascal, maxwell, debug, bigbox, mic, fermi)
Account	account billed for the job
Group	Group the person requesting the job belongs to
GID	Id for the group requesting the job
User	Name of user who ran the job
UID	Id of user who ran the job
Submit	time (UTC) that job was submitted
Eligible	time (UTC) when the job became eligible (promoted to the queue?)
Start	time (UTC) job started
End	time (UTC) job finished
Elapsed	time it took for job to run (is this ever different from End - Start?)
ExitCode	Code returned by job at exit(); the return code is followed by a colon and the code for the signal which triggered exit if the job was terminated by a signal
State	Job state or status (COMPLETED, CANCELLED, FAILED, TIMEOUT, PREEMPTED, etc.)
NNodes	Number of nodes in a job or step - if started, this will be actual number of nodes allocated, else it is the number requested
NCPUS	Count of Allocated CPUs
ReqCPUS	Required CPUs
ReqMem	Minimum memory for the job in MB
ReqGRES	Names and counts of Generic RESources requested
ReqTRES	Requested (minimum) Trackable RESource
Timelimit	Amount of time job should be allowed to run before it is aborted
NodeList	List of nodes in the job/step
JobName	Name of the job run by the Simple Linux Utility for Resource Management (SLURM)